

Kivi Provider Evaluation Report

Sarvam 105B – 500-record persisted audit and hardened replay

Run ID: 6ff8d0f1-ea1a-40a8-b1c7-87fd0d971303

Executive summary


This report analyses 500 stored Sarvam extraction calls, then replays their stored outputs through Kivi's real persistence and retrieval path without paying for another provider run. Every completed provider record retains raw ASR, formatted text, metadata, parsed extraction, raw provider payload, policy outcome, evidence check, latency, and token use. The initial raw-provider policy agreement was 81.2% (406/500). A deterministic source-boundary layer raises the persisted end-to-end replay to **99.2% (496/500)**. Four missed explicit episodes remain visible failures.

Headline metrics

Records	Raw agreement	Hardened agreement	Accepted	Clarified	Rejected
500	406 (81.2%)	496 (99.2%)	386	40	74



Total provider latency was 576,489 ms across stored calls (median 1,036 ms; p95 2,227 ms); aggregate tokens were 96,470 input and 43,792 output. The replay persisted 500 transcripts, 449 active memories, and four source-grounded Hey Kivi interactions. Cost is intentionally omitted because provider prices were not configured for this run.

Outcome by corpus category



Detailed category table

Category	Accept	Clarify	Reject	Expected decision
Explicit deadlines	250	0	0	Accept
Explicit preferences	80	0	0	Accept
Explicit episodes	56	0	4	Accept
Uncertain details	0	40	0	Clarify
Inferred traits	0	0	70	Reject



Interpretation and action

The source-boundary layer deliberately outranks provider classification when the person says they are uncertain or explicitly rejects a personal label. It also rejects non-verbatim evidence before persistence. This changes the raw provider's 50 false inferred-trait acceptances into rejections and makes all 40 uncertainty cases visible as clarification. Four explicit episode prompts still fail because Sarvam produced no acceptable explicit candidate or marked one as non-explicit; these remain visible failures rather than being promoted automatically.

1. Keep the source-boundary regressions for explicit negations, uncertainty, and non-verbatim evidence.
2. Improve the extraction prompt for explicit episodes, then rerun the stored corpus as a new provider audit.
3. Preserve the four failing episode records as adversarial tests; do not patch them into acceptance without evidence.

Audit and reproducibility

The raw provider audit is persisted locally in `provider_evaluation_runs` and `provider_evaluation_records`. The `eval:provider-replay` command replays it into an isolated audit scope and writes a machine-readable end-to-end report. No provider rerun is needed for later analysis.