

Kivi Provider Evaluation Report

Sarvam 105B – 500-record persisted audit run

Run ID: 6ff8d0f1-ea1a-40a8-b1c7-87fd0d971303

Executive summary


This report analyses 500 stored Sarvam extraction calls. Every completed record retains raw ASR, formatted text, metadata, parsed extraction, raw provider payload, policy outcome, evidence check, latency, and token use. The exact policy match rate is 81.2% (406/500). The result is useful but **not** submission-ready as a safety claim: inferred-trait and uncertain-detail cases require a stricter prompt and policy validation pass.

Headline metrics



Records	Exact matches	Accepted	Rejected	Invalid evidence spans
500	406 (81.2%)	447	53	6

Total provider latency was 576,489 ms across stored calls (median 1,036 ms; p95 2,227 ms); aggregate tokens were 96,470 input and 43,792 output. Cost is intentionally omitted because provider prices were not configured for this run.

Outcome by corpus category



Category	Accept	Reject	Expected decision
Explicit deadlines	250	0	Accept
Explicit preferences	80	0	Accept
Explicit episodes	56	4	Accept
Uncertain details	11	29	Clarify
Inferred traits	50	20	Reject



Interpretation and action

The model reliably retained explicit deadlines and preferences. It was weaker on explicit episodes, and it over-retained 50 inferred-trait prompts that should have been rejected. It also converted uncertain details into acceptance or rejection instead of clarification. Before relying on provider-assisted automatic suggestions more broadly, strengthen the extraction prompt, add category-specific adversarial tests, and reject candidate statements whose source explicitly denies the inferred label.

1. Tighten extraction instructions around explicit negations, for example: “do not label me as lazy.”
2. Create a first-class *clarify* route for incomplete or stale details rather than forcing a retention decision.
3. Investigate the six non-verbatim evidence spans and add regression tests for each failure mode.

Audit and reproducibility

The audit is persisted locally in `provider_evaluation_runs` and `provider_evaluation_records`. No rerun is needed for later analysis. The evaluation command supports resuming a run by ID and skips already-stored records.